

Neural Consensus Protocol (NCP)

Mathematical Foundations and Algorithm Specification

Version 1.0 — May 2026 | NCP Working Group

1. Formal Definitions

Definition 1 (Reasoning Chain). A reasoning chain C is a tuple (id, a, v, S, σ) where: $id \in H$ is a unique identifier, $a \in A$ is the owning agent, $v \in N$ is the version, $S = (s_1, s_2, \dots, s_n)$ is an ordered sequence of steps, $\sigma \in \{\text{active, broken, archived}\}$ is the chain status.

Definition 2 (Step). A step s_i is a tuple (i, c_i, τ_i, n_i) where: $i \in N$ is the index, c_i is the assertion content, $\tau_i \in \{\text{active, collapsed, fortified}\}$ is the status, $n_i \in N$ is the attack count.

Definition 3 (Attack). An attack X is a tuple $(a_src, C_tgt, i_tgt, r, \omega)$ where: a_src is the attacker, C_tgt is the target chain, i_tgt is the target step, r is the reason, $\omega \in \{0,1\}$ is the outcome.

Definition 4 (Address). Every step has a unique address:

$addr(s) = agent.v\{version\}.Step\{index\}$ e.g., GPT-5.4.v2.Step3

2. Entropy Calculations

2.1 10-Dimensional Entropy Vector

The system state is characterized by:

$H(t) = (H_{sem}, H_{cau}, H_{bnd}, H_{tmp}, H_{dep}, H_{div}, H_{inf}, H_{prp}, H_{evi}, H_{imp})$

where each $H_i \in [0, 1]$ measures a distinct type of uncertainty.

2.2 Divergence Entropy (Stance-Based)

For each active chain, define a stance vector $v_a \in \{-1, 0, +1\}^d$:

$v_{\{a,j\}} = +1$ if positive signals dominate, -1 if negative, 0 otherwise

Stance agreement between agents a, b :

$agree(a, b) = (1/|D|) \times \sum_{j \in D} match(v_{\{a,j\}}, v_{\{b,j\}})$

where $match = 1$ if equal, 0.5 if one is 0 , 0 otherwise.

Divergence entropy:

$H_{div} = 1 - (2 / |A|(|A|-1)) \times \sum_{a < b} agree(a, b)$

2.3 Total Entropy with Impact Amplification

Theorem 1. The total entropy is amplified by impact:

$$H_{\text{total}} = \min(1, H_{\text{mean}} \times (1 + 0.5 \times H_{\text{imp}}))$$

Amplification range: 1.0× (no impact) to 1.5× (max impact).

3. UCB1 Attack Scheduling

3.1 Multi-Armed Bandit Formulation

Definition 5 (UCB1 Score). For step s :

$$\text{UCB}(s) = c/n + \sqrt{2 \times \sqrt{(\ln(N)/n)} + 0.1 \times \ln(1 + \Delta t)}$$

where c = collapse count, n = attack count, N = total attacks, Δt = rounds since last attack

Property 1 (Exploration Guarantee). If $n = 0$, then $\text{UCB}(s) = +\infty$. Every step is attacked at least once before any step is attacked twice.

Theorem 2 (Coverage Bound). With $|A|$ agents and $|S|$ total steps, after r rounds:

$$\text{coverage}(r) \geq \min(1, r \times |A| / |S|)$$

With 30 agents and 210 steps: full coverage in $\lceil 210/30 \rceil = 7$ rounds.

4. Decentralized Parallel Voting

Definition 6. Voting set: $V = A \setminus \{a_{\text{src}}\}$ (all except attacker)

Each voter independently: $\text{vote}(v) \in \{\text{collapsed}, \text{defended}\}$

Majority Decision:

$$\omega = 1 \text{ (collapsed) if } |\{v: \text{vote}(v)=\text{collapsed}\}| > |\{v: \text{vote}(v)=\text{defended}\}|, \text{ else } 0$$

Property 2 (No Privileged Node). No judge, no arbiter. Attacker cannot vote.

Property 3 (Parallel Execution). $\text{latency}(\text{voting}) = \max_{v \in V} \text{latency}(v) = O(1)$ w.r.t. $|V|$

Theorem 3 (Byzantine Fault Tolerance). With n voters, tolerates $\lfloor (n-1)/2 \rfloor$ Byzantine nodes.

Proof: Byzantine voter contributes at most 1 wrong vote. Honest majority $n-f > n/2$ when $f < n/2$.

5. Convergence Theory

Definition 7 (Verified Convergence). System converges when:

$$\text{coverage}(t) > \theta_c \text{ AND } \text{verified_ratio}(t) > \theta_v \text{ AND } t \geq t_{\text{min}}$$

Default: $\theta_c = 0.5$, $\theta_v = 0.4$, $t_{\text{min}} = 2$

Theorem 4 (Convergence Speed). Expected rounds:

$$E[T] \leq \lceil \theta_c \times |S| / |A| \rceil + 1$$

Corollary: 30 agents, 49 steps, $\theta_c=0.5$: $E[T] \leq 2$ rounds.

This proves: **more agents** → **faster convergence**.

6. ELO Reputation System

After attack between attacker a and defender d:

$$E_a = 1 / (1 + 10^{((R_d - R_a)/400)})$$

$$R_a' = R_a + K \times (\omega - E_a)$$

$$R_d' = R_d + K \times ((1-\omega) - E_d)$$

$K = 32$, $\omega = 1$ if attack succeeds, 0 otherwise.

Vote weight mapping:

$$w(a) = \text{clip}((R_a - 800)/400 + 0.5, 0.5, 2.0)$$

7. Consensus Persistence with Decay

Confidence decay function:

$$\phi(t) = \max(\phi_{\text{floor}}, 1.0 - \lambda(t-t_{\text{ref}}) + \mu \times n_{\text{ref}}(t))$$

$\lambda = 0.005/\text{day}$, $\phi_{\text{floor}} = 0.3$, $\mu = 0.02/\text{reference}$

$$\text{Challenge penalty: } \phi' = \phi - 0.15$$

Revocation propagation: if $\phi' \leq 0.1$, all dependents lose 0.2

7.1 Cross-Session α Calculation

$$\alpha(c) = \sum_{a \in A} w(a) \times 1[c \in C_a] / \sum_{a \in A} w(a)$$

Convergence threshold: $\alpha \geq 0.65 \rightarrow \text{TRUE_CONSENSUS}$

8. Protocol Invariants

- I1: $\forall C: v_C \geq 1$ (version is positive)
- I2: $\forall s: n_s \geq 0$ (attack count non-negative)
- I3: $\tau_s = \text{fortified} \Rightarrow n_s \geq 2$ (fortification requires 2+ survived attacks)
- I4: $\sigma_C = \text{broken} \Rightarrow \exists s \in C.S: \tau_s = \text{collapsed}$
- I5: $\text{REPAIR}(C) \Rightarrow v_{\{C\}} = v_C + 1$ (repair increments version)
- I6: $\sigma_C = \text{archived} \Rightarrow C$ cannot be attacked
- I7: $\text{vote}(a_{\text{src}}, X) = \perp$ (attacker cannot vote on own attack)
- I8: $H_{\text{total}} \in [0, 1]$ (entropy is bounded)
- I9: $\alpha(c) \in [0, 1]$ (consensus degree is bounded)
- I10: $\phi(t) \geq \phi_{\text{floor}}$ unless revoked

9. Complexity Summary

Component	Per-Round Complexity	Bottleneck
Chain parsing	$O(A \times k)$	Regex
UCB1 scoring	$O(S)$	Step traversal
Attack routing	$O(A \times S)$	Step build
Parallel voting	$O(1)$ latency	Slowest model
Entropy	$O(A \times k)$	Stance extraction
α extraction	$O(1)$ call	Single model
Consensus store	$O(C_store)$	Keyword match

True bottleneck: LLM inference latency (~10-30s). All algorithmic overhead is sub-second.

10. Scaling Law

Agents	Attacks/Round	Coverage/Round	Expected Convergence
3	3	~15%	4-6 rounds
7	7	~35%	2-3 rounds
15	15	~70%	1-2 rounds
30	30	~95%	1 round

The fundamental scaling law: More agents → more parallel attacks → higher coverage per round → faster convergence. This is the opposite of traditional multi-agent systems.